





3D Object Recognition and 6DOF Pose Estimation

Aitor Aldoma

May 10, 2013

- ▶ This morning, we saw "How does a good feature look like?"
 - ▶ Representing shapes (global, local) in a compact manner.
 - ▶ How to match features to define **correspondences** between source and target.

- ▶ This morning, we saw "How does a good feature look like?"
 - ▶ Representing shapes (global, local) in a compact manner.
 - ▶ How to match features to define **correspondences** between source and target.
- ▶ In this talk:
 - ▶ Given a set of models (training data), how do we recognize them and estimate their pose in a particular scene?





- ▶ `pcl::keypoints`, `pcl::features` covered previously in the morning.
- ▶ In this talk, we focus on `CorrespondenceGrouping` and `Hypothesis Verification`.
- ▶ In contrast to registration, we simulatenously deal with several models.
- ▶ Other options in PCL:
 - ▶ LINEMOD [HinterstoisserPAMI2012]
 - ▶ ORR [PapazovACCV2010]
 - ▶ Segmentation + global features

1. Correspondence grouping

- ▶ Given a set of correspondences (models - scene), group them together in geometrically consistent clusters from which the pose of the models can be extracted.
- ▶ In contrast to RANSAC based methods, allows simultaneous recognition of multiple objects.
- ▶ Usually applied on recognition pipelines based on local features.

1. Correspondence grouping

- ▶ Given a set of correspondences (models - scene), group them together in geometrically consistent clusters from which the pose of the models can be extracted.
- ▶ In contrast to RANSAC based methods, allows simultaneous recognition of multiple objects.
- ▶ Usually applied on recognition pipelines based on local features.

2. Hypothesis Verification; given a set of object hypotheses with a 6DoF pose.

- ▶ Aims at removing false positives while keeping true positives.
- ▶ Allows merging hypotheses coming from different pipelines in a principled way.

1. Correspondence grouping

- ▶ Given a set of correspondences (models - scene), group them together in geometrically consistent clusters from which the pose of the models can be extracted.
- ▶ In contrast to RANSAC based methods, allows simultaneous recognition of multiple objects.
- ▶ Usually applied on recognition pipelines based on local features.


2. Hypothesis Verification; given a set of object hypotheses with a 6DoF pose.

- ▶ Aims at removing false positives while keeping true positives.
- ▶ Allows merging hypotheses coming from different pipelines in a principled way.

PCL module: `pcl::recognition`

From 3D models to 2.5D data

- Simulate input from depth/3D sensors.



```

1  typedef pcl::PointCloud<pcl::PointXYZ>::Ptr CloudPtr;
2  pcl::apps::RenderViewsTessellatedSphere render_views;
3  render_views.setResolution (resolution_);
4  render_views.setTessellationLevel (1); //80 views
5  render_views.addModelFromPolyData (mapper); //vtk model
6  render_views.setGenOrganized(false);
7  render_views.generateViews ();
8  std::vector< CloudPtr > views;
9  std::vector< Eigen::Matrix4f > poses;
10 render_views.getViews (views);
11 render_views.getPoses (poses);


```


1. Introduction
2. Correspondence Grouping
3. Hypothesis Verification
4. A recognition example


- Incrementally build clusters of correspondences that are geometrically consistent:

$$|||p_{i,m} - p_{j,m}||_2 - ||p_{i,s} - p_{j,s}||_2| < \varepsilon \quad (1)$$

- All elements in cluster are geom. consistent to each other.
- Parameters:
 - ε : keypoint inaccuracy, noise
 - gc_min_size : minimum cluster size (at least 3)




- ▶ GC constraint is weak, specially for small consensus sizes!
- ▶ 6D space projected to 1D!
- ▶ Should be specially used when data is noisy or presents artifacts that do not permit to compute a repeatable RF (see Hough3D).



- ▶ How to use it within PCL?
- ▶ `m_s_corrs` are correspondences with indices to `m_keypoints` and `s_keypoints`.

```
1  pcl::CorrespondencesPtr m_s_corrs; //fill it
2  std::vector<pcl::Correspondences> clusters;
3  pcl::GeometricConsistencyGrouping<PT, PT> gc_clusterer;
4  gc_clusterer.setGCSize (cg_size);
5  gc_clusterer.setGCThreshold (cg_thres);
6  gc_clusterer.setInputCloud (m_keypoints);
7  gc_clusterer.setSceneCloud (s_keypoints);
8  gc_clusterer.setModelSceneCorrespondences (m_s_corrs);
9  gc_clusterer.cluster (clusters);
```

- ▶ Correspondence votes are accumulated in a 3D Hough space. [TombarilPSJ2012]
- ▶ Each point associated with a repeatable RF, RFs used to:
 - ▶ reduce voting space from 6 to 3D...
 - ▶ ... by reorienting the voting location
- ▶ Local maxima in the Hough space identify object instances (handles the presence of multiple instances of the same model in the scene)



- ▶ How to use it within PCL?
- ▶ `m_s_corrs` are correspondences with indices to `m_keypoints` and `s_keypoints`.


```
1  typedef pcl::ReferenceFrame RFTYPE;  
2  pcl::PointCloud<RFTYPE>::Ptr model_rf; //fill with RFs  
3  pcl::PointCloud<RFTYPE>::Ptr scene_rf; //fill with RFs  
4  pcl::CorrespondencesPtr m_s_corrs; //fill it  
5  std::vector<pcl::Correspondences> clusters;  
6  
7  pcl::Hough3DGrouping<PT, PT, RFTYPE, RFTYPE> hc;  
8  hc.setHoughBinSize (cg_size);  
9  hc.setHoughThreshold (cg_thres);  
10 hc.setUseInterpolation (true);  
11 hc.setUseDistanceWeight (false);  
12 hc.setInputCloud (m_keypoints);  
13 hc.setInputRf (model_rf);  
14 hc.setSceneCloud (s_keypoints);  
15 hc.setSceneRf (scene_rf);  
16 hc.setModelSceneCorrespondences (m_s_corrs);  
17 hc.cluster (clusters);
```

1. Introduction
2. Correspondence Grouping
- 3. Hypothesis Verification**
4. A recognition example

- ▶ Keep along the recognition pipeline as many hypotheses as possible and use HV to select those best "explaining the scene".
- ▶ A hypothesis \mathcal{M}_i is a model aligned to the scene \mathcal{S} .
- ▶ Main **goal**: Remove FPs without rejecting TPs.



- ▶ Keep along the recognition pipeline as many hypotheses as possible and use HV to select those best "explaining the scene".
- ▶ A hypothesis \mathcal{M}_i is a model aligned to the scene \mathcal{S} .
- ▶ Main **goal**: Remove FPs without rejecting TPs.
- ▶ 3 options in PCL:
 - ▶ **Greedy** [AldomaDAGM12]
 - ▶ **Conflict Graph** [PapazovACCV11]
 - ▶ **Global HV** [AldomaECCV12]



- ▶ Reasoning about occlusions to handle occluded objects (in common with all 3 methods).
- ▶ For each hypothesis \mathcal{M}_i , count $\#inliers$ and $\#outliers$.
- ▶ Greedily select the best hypothesis ($\#inliers - \lambda \cdot \#outliers$) ...
- ▶ ... and update the inliers count for successive hypotheses, resort and repeat.
- ▶ \mathcal{M}_i selected if $\#inliers - \lambda \cdot \#outliers > 0$.

```
1 pcl::GreedyVerification<pcl::PointXYZ, pcl::PointXYZ>  
   greedy_hv(lambda);  
2 greedy_hv.setResolution (0.005f);  
3 greedy_hv.setInlierThreshold (0.005f);  
4 greedy_hv.setSceneCloud (scene);  
5 greedy_hv.addModels (aligned_hypotheses, true);  
6 greedy_hv.verify ();  
7 std::vector<bool> mask_hv;  
8 greedy_hv.getMask (mask_hv);
```

- ▶ First, a sequential stage that discards hypotheses based on percentage of inliers and outliers.
- ▶ From the remaining hypotheses, some are selected based on a non-maxima suppression stage on a conflict graph.
- ▶ Two hypothesis are in conflict if they share the same space.


```
1 pcl::PapazovHV<pcl::PointXYZ, pcl::PointXYZ> papazov;
2 papazov.setResolution (0.005f);
3 papazov.setInlierThreshold (0.005f);
4 papazov.setSupportThreshold (0.08f); //inliers
5 papazov.setPenaltyThreshold (0.05f); //outliers
6 papazov.setConflictThreshold (0.02f);
7 papazov.setSceneCloud (scene);
8 papazov.addModels (aligned_hypotheses, true);
9 papazov.verify ();
10 std::vector<bool> mask_hv;
11 papazov.getMask (mask_hv);
```

- ▶ Consider the two possible states of a single hypothesis $x_i = \{0, 1\}$ (inactive/active).
- ▶ By switching the state of an hypothesis, we can evaluate a global cost function that tell us how good the current solution $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ is.
- ▶ Formally, we are looking for a solution $\tilde{\mathcal{X}}$ such that:

$$\tilde{\mathcal{X}} = \underset{\mathcal{X} \in \mathbb{B}^n}{\operatorname{argmin}} \{ \mathfrak{F}(\mathcal{X}) = f_{\mathcal{S}}(\mathcal{X}) + \lambda \cdot f_{\mathcal{M}}(\mathcal{X}) \} \quad (2)$$

- ▶ $\mathfrak{F}(\mathcal{X})$ considers the whole set of hypotheses (\mathcal{M}) as a global scene model instead of considering each model hypothesis separately.


- ▶ $\mathfrak{F}(\mathcal{X})$ simultaneously enforces cues defined on the scene \mathcal{S} as well as cues defined on the set of hypothesis, \mathcal{M} .



- ▶ Given a certain configuration of $\mathcal{X} = \{x_1, \dots, x_n\}$:
 - ▶ Maximize number of scene points explained (orange).
 - ▶ Minimize number of model outliers (green).
 - ▶ Minimize number of scene points multiple explained (black).
 - ▶ Minimize number of unexplained scene points close to active hypotheses (yellow, purple)
- ▶ Optimization solved using Simulated Annealing or other metaheuristics (METSlib library).

```
1  pcl::GlobalHypothesesVerification<pcl::PointXYZ, pcl::
    PointXYZ> go;
2  go.setResolution (0.005f);
3  go.setInlierThreshold (0.005f);
4  go.setRadiusClutter (0.04f);
5  go.setRegularizer (3.f); //outliers' model weight
6  go.setClutterRegularizer (5.f); //clutter points weight
7  go.setDetectClutter (true);
8  go.setSceneCloud (scene);
9  go.addModels (aligned_hypotheses, true);
10 go.verify ();
11 std::vector<bool> mask_hv;
12 go.getMask (mask_hv);
```

1. Introduction
2. Correspondence Grouping
3. Hypothesis Verification
4. A recognition example



► 5 object models.



/home/aitor/data/Mars_dataset/scenes/pcl_scenes/t07.ply.pcd



Recognition results



Ground truth